


K Nearest Neighbors

In K-Nearest Neighbors (KNN), the class of a new data point is predicted based on the majority class among its k closest training samples, measured using a distance metric like Euclidean distance.

22 OCT, 2025



BAYESIANBRAT*



Nearest Neighbors (NNs)

- It is a simple supervised learning model.
- Given an instance x_0 , determine nearest neighbor = instance in training data closest to x_0
- Assign x_0 to the class of the nearest neighbor
- Often use K nearest neighbors (NNs)
- Assign x_0 to the majority class of the K NNs

Euclidean distance is measured as

$$d(x, y) = \sqrt{\sum (x_i - y_i)^2}$$

- K is a hyperparameter: a parameter set ~~by~~ before the model training process begins.
- Hyperparameter = a setting you choose before training.
- Not learned from data
- Instead you decide its value and test which works best.

For K -NN

K = how many neighbors you look at

small $K \rightarrow$ very flexible, but noisy

large $K \rightarrow$ smoother, but might miss local structure

For $K=1$ or any small K , causes overfitting

For K = Big, there's a risk of underfitting

Also, K NN is a ^{almost} ~~almost~~ lazy learner because it ^{almost} learns nothing during the training stage. Because during prediction, it recalculates everything.

\rightarrow Euclidean distance:

$$d(x, y) = \sqrt{\sum (x_i - y_i)^2}$$

\rightarrow The straight-line distance (Pythagoras)

\rightarrow Manhattan distance:

$$d(x, y) = \sum |x_i - y_i|$$

\rightarrow Minkowski distance:

→ Minkowski Distance:

This is a general formula that covers both Euclidean and Manhattan

$$d(x, y) = \left(\sum_{i=1}^n |x_i - y_i|^p \right)^{1/p}$$

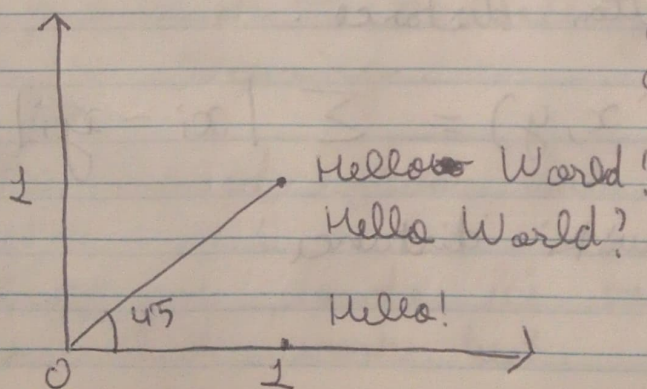
★ Cosine Similarity:

→ Cosine similarity is determined by the $\cos(\theta)$, where θ is the angle between two vectors.

→ These two vectors are formed by a point on graph to origin and some another point to the origin of the graph.

⇒ Example:

Words	Hello	World
Hello World!	1	1
Hello!	1	0
Hello World?	1	1



$$\cos(\theta) = 1$$
$$\cos(45) = 0.71$$

There's a generalized formula:

$$\text{Cosine similarity} = \frac{\sum_{i=1}^n A_i B_i}{\sqrt{\sum_{i=1}^n A_i^2} \sqrt{\sum_{i=1}^n B_i^2}}$$

Exercise:

	I love Troll 2	Gymkhana
A: I love Troll 2!	1 1 1 1 0	0
B: I love Gymkhana!	1 1 0 0 1	1

Applying formula:

$$= \frac{(1 \times 1) + (1 \times 1) + (1 \times 0) + (1 \times 0) + (0 \times 1)}{\sqrt{1^2 + 1^2 + 1^2 + 1^2 + 0^2} \sqrt{1^2 + 1^2 + 0^2 + 0^2 + 1^2}}$$

=

2

$$\sqrt{4} \sqrt{3}$$

=

1

$$\sqrt{3}$$

$$= 0.58$$

→ In summary, the cosine similarity is a relatively easy to calculate metric that tells us how similar or different things are.

$$\boxed{\text{Cosine distance} = 1 - \text{Cosine Similarity}}$$

★ Categorical Attributes:

- So far, we used numeric distances (Euclidean, Manhattan etc.)
- But what if attributes like ("Wet/dry", "smooth/rough") need to be calculated?
- We can't subtract them, instead we use match distance

The example in the PDF
We have a query instance:

$x_0 = (\text{Weather} = \text{Wet}, \text{surface} = \text{medium}, \text{experience} = \text{low}, \text{grade} = \text{flat})$

instance	Weather	surface	experience	grade
0	dry	smooth	high	steep
1	dry	smooth	high	flat
2	dry	smooth	high	flat
3	wet	rough	low	steep
4	dry	rough	low	steep

→ Take instance 0 vs x_0 :

Weather: Wet vs dry → different → 1
 surface: mid vs smooth → different → 1
 experience: low vs high → different → 1
 grade: flat vs steep → different → 1

Total distance $\rightarrow 4$

Take instance 3 vs x_0 :

Weather: Wet vs Wet $\rightarrow 0$

Surface: mid vs rough $\rightarrow 1$

exp: low vs low $\rightarrow 0$

grade: flat vs steep $\rightarrow 1$

Total distance $= 2$

Likewise we calculate for all other instances.

★ Attribute Scales

- \rightarrow Many distance measures depends on the scale of attribute (unit of measure)
- \rightarrow Often important to scale the attributes to unitless quantities.

Standardization formula:

$$x^* = \frac{x - \text{mean}(x)}{\text{stddev}(x)}$$

Example

x_1	x_2	x_1^*	x_2^*
1.8	3.7	-0.2641	0.48668
2.0	3.0	-0.15288	0.27762
2.5	3.5	0.12523	0.42695
2.0	4.1	-0.15288	0.6015
1.8	1.0	-0.26413	-0.31931
2.3	3.8	0.01398	0.51653
4.0	6.0	0.9595	
-1.0	3.3	-1.8215	
6.0	-4.0	2.0720	
1.3	-3.7	-0.5125	

How did we calculate

$$x_1 \text{ mean} = 2.3 \quad \text{stddev} = 1.798$$

$$x_2 \text{ mean} = 2.1 \quad \text{stddev} = 3.348$$

$$x_1^* = \frac{1.8 - 2.3}{1.7978} = -0.264$$

$$x_2^* = \frac{3.7 - 2.1}{3.348} = 0.487$$

Weighted KNN

- In standard KNN, you pick the k nearest neighbors.
- Each neighbor gets equal weight → simple major vote (classification) or average (regression).

Problem:

- What if one neighbor is much closer to x_0 than others?
- Shouldn't it have more influence?

Weighted KNN

- Instead of giving all neighbors equal say, give closer neighbors higher weight.

$$\hat{y}(x_0) = \frac{\sum_{i=1}^k w_i y_i}{\sum_{i=1}^k w_i}$$

$$\sum_{i=1}^k w_i$$

Where

$$w_i = \frac{1}{d(x_0, x_i)}$$

sometimes $w_i = \frac{1}{d(x_0, x_i)^2}$ etc.

So, the nearest neighbors vote counts more than a farther one.

→ Example (Classification)

Red at 0.5

Blue at 0.6

Red at 1.2

Unweighted K-NN (K=3)

Red = 2 votes, Blue = 1 vote → Predicted Red

Weighted K-NN ($1/d$ rule)

Red → 0.5 → $\frac{1}{0.5} = 2.0$

Blue → 0.6 → $\frac{1}{0.6} = 1.67$

Red → 1.2 → $\frac{1}{1.2} = 0.83$

Total weights = Red : 2.83
Blue : 1.67

Answer: Red

~~Example~~

Example data (from before)

	dist	weight
Red	0.5	2.0
Blue	0.6	1.67
Red	1.2	0.83

$$\text{Total Weight} = 2.0 + 1.6667 + 0.8333 \\ = 4.5$$

Classification (Weighted vote):

$$\begin{aligned} \text{Red } w_1 + w_3 &= 2.0 + 0.8333 = 2.8333 \\ w_2 &= 1.6667 \end{aligned}$$

Regression (if labels are numeric, Red=1, Blue=0)

$$\begin{aligned} \hat{y} &= \frac{2 \times 1 + 1.6667 \cdot 0 + 0.8333 \times 1}{4.5} \\ &= 0.6296 \end{aligned}$$